


A small task-score gain. A real token trade-off.

Held-out synthetic user histories · 2 personas × 8 questions × 2 repeats · 32 trials per condition

Exact-task success


Higher is better



+12.5 percentage points
Wiki vs full history · exact-task success

Foreground tokens


Thousands · lower is better · 32 trials total



24.7% fewer tokens
Wiki vs notes + sessions · foreground only

Median answer latency

Lower is better · seconds



37.0% more tokens
Wiki vs full history · foreground only

WHAT THIS SUPPORTS

A narrow observed improvement in task completion, not established broad semantic gains or cost savings.

Exact JSON / expected-value scoring; execution and formatting affect results. All failed trials remain included.
Shared task templates, two held-out personas; no confidence intervals. Snapshot preparation is outside latency.
Foreground charts exclude compilation: 402,516 recorded tokens across development + held-out compilation attempts, including estimated embeddings. That total is not a per-condition charge. No monetary cost comparison.